

2- Feature Selection

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: %matplotlib inline

import os, sys
from time import time

from pathlib import Path
import numpy as np
import pandas as pd
import pandas_datareader.data as web

import statsmodels.api as sm
from sklearn.feature_selection import mutual_info_regression
from sklearn.preprocessing import scale
import lightgbm as lgb
from scipy.stats import spearmanr
import shap

from alphasens.tears import (create_returns_tear_sheet,
                             create_summary_tear_sheet,
                             create_full_tear_sheet)

from alphasens import plotting
from alphasens import performance as perf
from alphasens import utils

import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: sns.set_style('whitegrid')
idx = pd.IndexSlice
deciles = np.arange(.1, 1, .1).round(1)
```

```
In [4]: MONTH = 21  
        YEAR = 252
```

Load Data

```
In [ ]: DATA_STORE = 'data/stock_prices.h5'
```

```
In [ ]: factors = (pd.read_hdf(DATA_STORE, 'model_data').sort_index())
```

```
In [7]: factors.info(show_counts=True)
```

```

<class 'pandas.core.frame.DataFrame'>
MultiIndex: 1403584 entries, ('A', Timestamp('2006-01-03 00:00:00')) to ('ZMH', Timestamp('2015-07-02 00:00:00'))
Data columns (total 55 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ex-dividend     1403584 non-null float64
1   split_ratio     1403584 non-null float64
2   adj_open        1403584 non-null float64
3   adj_high        1403584 non-null float64
4   adj_low         1403584 non-null float64
5   adj_close       1403584 non-null float64
6   adj_volume      1403584 non-null float64
7   ret_01          1403084 non-null float64
8   ret_03          1402084 non-null float64
9   ret_05          1401084 non-null float64
10  ret_10          1398584 non-null float64
11  ret_21          1393084 non-null float64
12  ret_42          1382584 non-null float64
13  ret_63          1372084 non-null float64
14  ret_126         1340584 non-null float64
15  ret_252         1277584 non-null float64
16  ret_fwd         1403584 non-null float64
17  BB_UP           1394084 non-null float64
18  BB_LOW          1394084 non-null float64
19  BB_SQUEEZE      1394084 non-null float64
20  HT              1372084 non-null float64
21  SAR             1403084 non-null float64
22  ADX             1390084 non-null float64
23  ADXR           1383584 non-null float64
24  PPO            1391084 non-null float64
25  AARONOSC        1396584 non-null float64
26  BOP             1403584 non-null float64
27  CCI            1397084 non-null float64
28  MACD           1387084 non-null float64
29  MACD_SIGNAL     1387084 non-null float64
30  MACD_HIST       1387084 non-null float64
31  MFI            1396584 non-null float64
32  RSI            1396584 non-null float64
33  STOCHRSI        1389084 non-null float64
34  STOCH           1395066 non-null float64

```

35	ULTOSC	1389584	non-null	float64
36	WILLR	1397084	non-null	float64
37	AD	1403584	non-null	float64
38	ADOSC	1396996	non-null	float64
39	OBV	1403583	non-null	float64
40	ATR	1396584	non-null	float64
41	ALPHA_63	1334328	non-null	float64
42	MARKET_63	1334328	non-null	float64
43	SMB_63	1334328	non-null	float64
44	HML_63	1334328	non-null	float64
45	RMW_63	1334328	non-null	float64
46	CMA_63	1334328	non-null	float64
47	ALPHA_252	1239828	non-null	float64
48	MARKET_252	1239828	non-null	float64
49	SMB_252	1239828	non-null	float64
50	HML_252	1239828	non-null	float64
51	RMW_252	1239828	non-null	float64
52	CMA_252	1239828	non-null	float64
53	month	1403584	non-null	uint8
54	weekday	1403584	non-null	uint8

dtypes: float64(53), uint8(2)
memory usage: 575.8+ MB

To get all features, we'll select the columns NOT containing forward returns:

```
In [8]: fwd_returns = factors.filter(like='fwd').columns
features = factors.columns.difference(fwd_returns).tolist()
```

So here are the available features:

```
In [9]: features
```

```
Out[9]: ['AARONOSC',  
        'AD',  
        'ADOSC',  
        'ADX',  
        'ADXR',  
        'ALPHA_252',  
        'ALPHA_63',  
        'ATR',  
        'BB_LOW',  
        'BB_SQUEEZE',  
        'BB_UP',  
        'BOP',  
        'CCI',  
        'CMA_252',  
        'CMA_63',  
        'HML_252',  
        'HML_63',  
        'HT',  
        'MACD',  
        'MACD_HIST',  
        'MACD_SIGNAL',  
        'MARKET_252',  
        'MARKET_63',  
        'MFI',  
        'OBV',  
        'PPO',  
        'RMW_252',  
        'RMW_63',  
        'RSI',  
        'SAR',  
        'SMB_252',  
        'SMB_63',  
        'STOCH',  
        'STOCHRSI',  
        'ULTOSC',  
        'WILLR',  
        'adj_close',  
        'adj_high',  
        'adj_low',  
        'adj_open',
```

```
'adj_volume',  
'ex-dividend',  
'month',  
'ret_01',  
'ret_03',  
'ret_05',  
'ret_10',  
'ret_126',  
'ret_21',  
'ret_252',  
'ret_42',  
'ret_63',  
'split_ratio',  
'weekday']
```

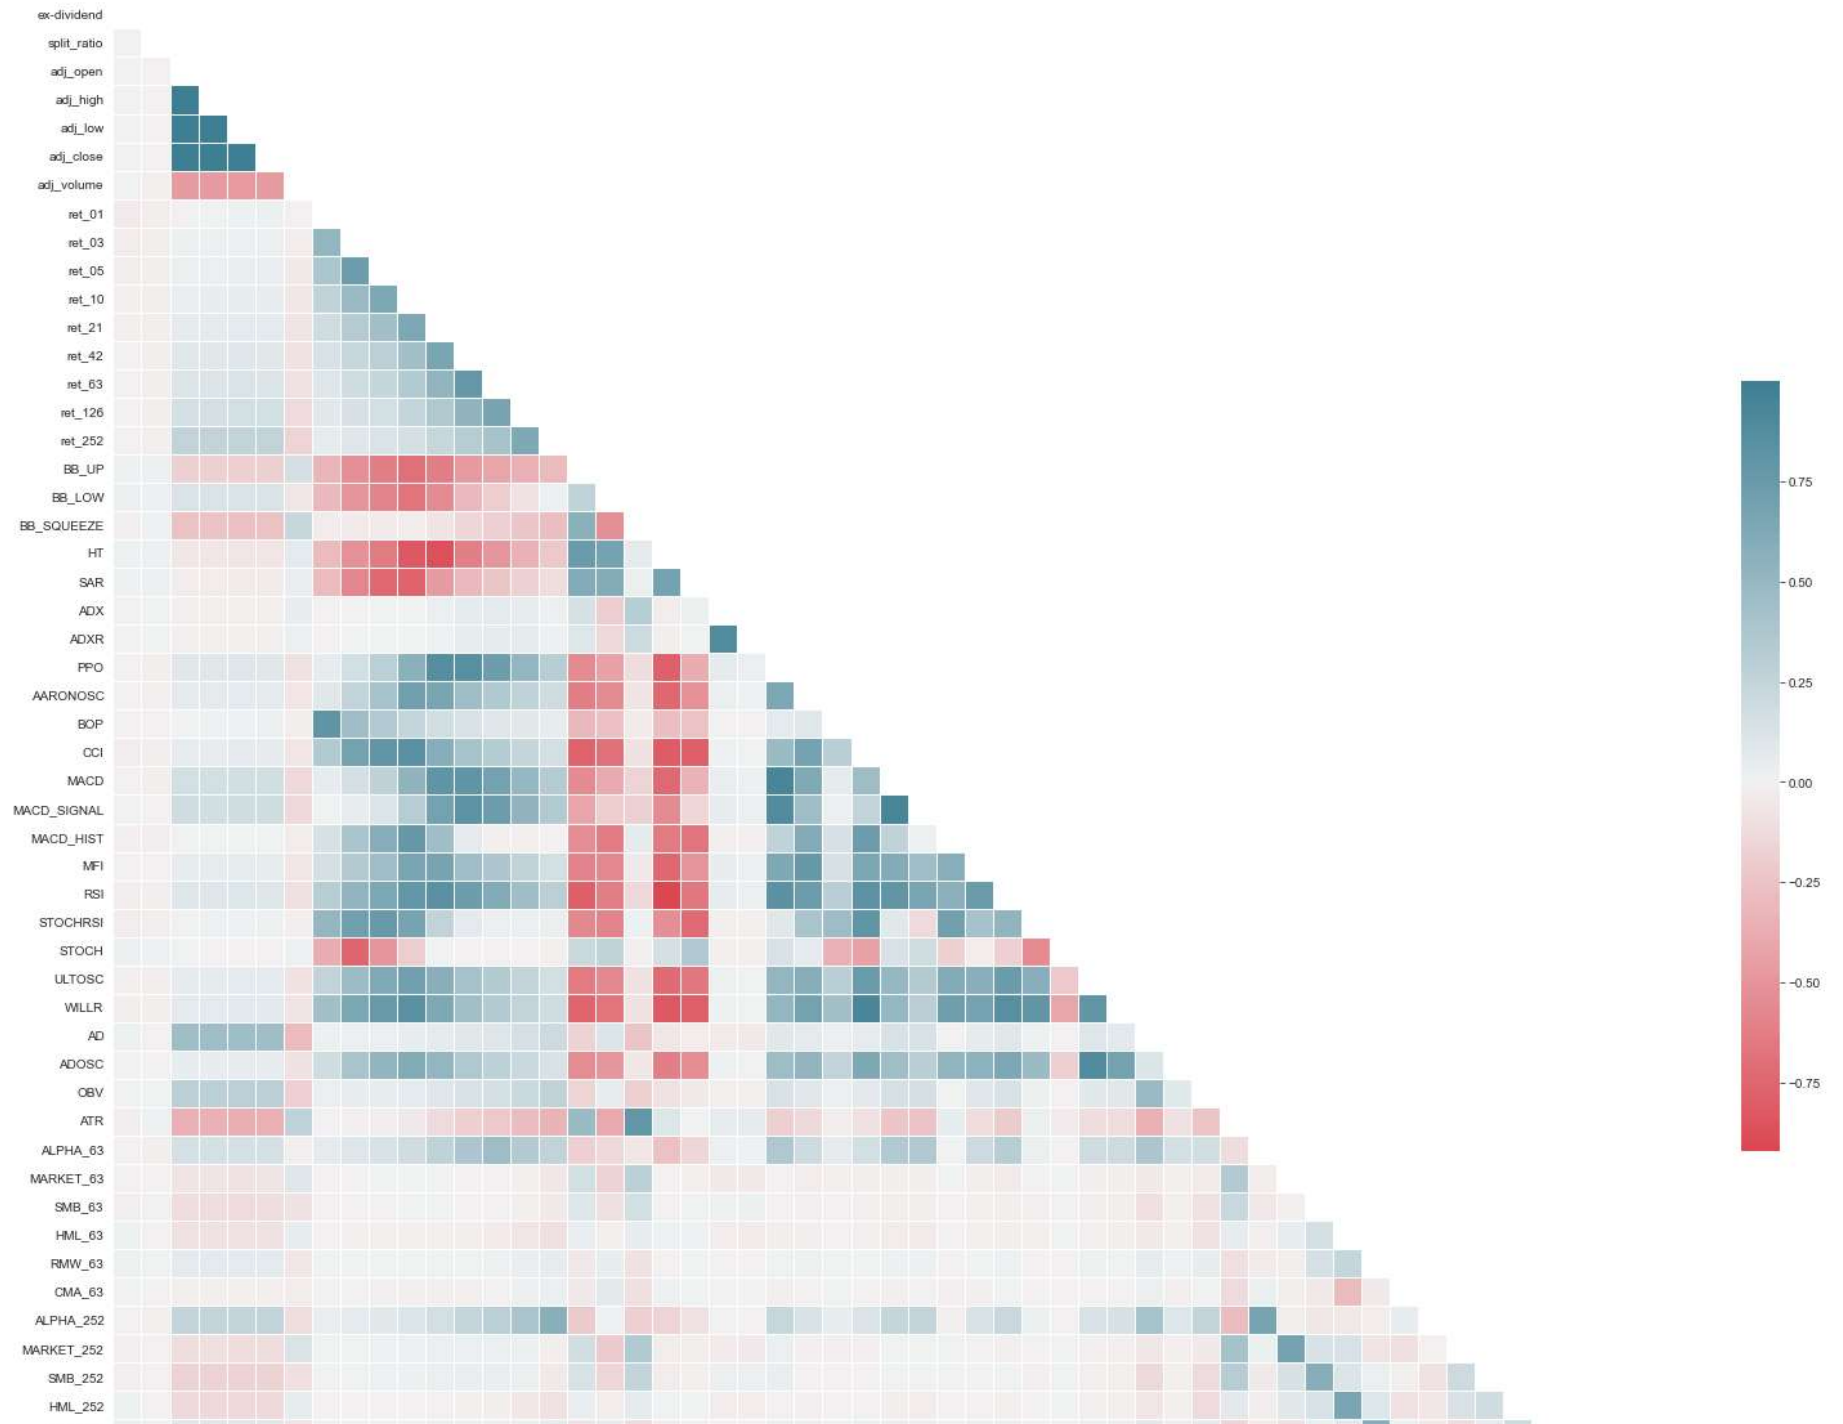
Factor Correlation


Which features are most alike in terms of their (rank) correlation?

```
In [10]: corr_common = factors.drop(fwd_returns, axis=1).corr(method='spearman')
```

Based on [this](#) example from the seaborn [docs](#).

```
In [11]: mask = np.triu(np.ones_like(corr_common, dtype=np.bool))  
fig, ax = plt.subplots(figsize=(22, 18))  
cmap = sns.diverging_palette(10, 220, as_cmap=True)  
  
sns.heatmap(corr_common, mask=mask, cmap=cmap, center=0,  
            square=True, linewidths=.5, cbar_kws={"shrink": .5})  
fig.tight_layout();
```






The clusters show a few clusters of features that are relatively highly correlated (in absolute terms). If you are interested in reducing the number of features to speed up training or reduce memory, selecting one representative of each cluster would be a reasonable approach.

```
In [12]: g = sns.clustermap(corr_common, cmap=cmap, figsize=(15, 15))
```




To inspect the correlation matrix, let's isolate the unique values from the correlation matrix, excluding the diagonal:

```
In [13]: corr_ = corr_common.stack().reset_index() # move column labels to rows as 2nd MultiIndex Level
corr_.columns = ['x1', 'x2', 'rho']
corr_ = corr_[corr_.x1 != corr_.x2].drop_duplicates('rho')
```

Top ten most correlated features

Select highest and lowest five correlation values:

```
In [14]: corr_.nlargest(5, columns='rho').append(corr_.nsmallest(5, columns='rho'))
```

Mutual Information


```
In [18]: mi = {}
for feature in features:
    print(feature)
    df = (factors
         .loc[:, ['ret_fwd', feature]]
         .dropna()
         .sample(n=100000)) # if it takes too long or you run into resource constraints, reduce the sample size
    discrete_features = df[feature].nunique() < 20
    mi[feature] = mutual_info_regression(X=df[[feature]],
                                       y=df.ret_fwd,
                                       discrete_features=discrete_features)[0]
```

AARONOSC
AD
ADOSC
ADX
ADXR
ALPHA_252
ALPHA_63
ATR
BB_LOW
BB_SQUEEZE
BB_UP
BOP
CCI
CMA_252
CMA_63
HML_252
HML_63
HT
MACD
MACD_HIST
MACD_SIGNAL
MARKET_252
MARKET_63
MFI
OBV
PPO
RMW_252
RMW_63
RSI
SAR
SMB_252
SMB_63
STOCH
STOCHRSI
ULTOSC
WILLR
adj_close
adj_high
adj_low
adj_open

```
adj_volume  
ex-dividend  
month  
ret_01  
ret_03  
ret_05  
ret_10  
ret_126  
ret_21  
ret_252  
ret_42  
ret_63  
split_ratio  
weekday
```

```
In [19]: mi = pd.Series(mi)  
mi.to_csv('mutual_info.csv')
```

```
In [ ]: mi.nlargest(25).sort_values().plot.barh(figsize=(14,8))  
sns.despine();
```



Comparison


```
In [ ]: top_ranked = stats.abs().rank(ascending=False).mean(1)
```

```
In [2]: #Note: feature based on mutual information did not used in the final tuning. Also features selection changed for the
```

```
In [ ]: top_ranked.to_csv('top_features.csv')
```



```
In [ ]: top_ranked.drop(categoricals).nsmallest(20).plot.bar(figsize=(16, 4), rot=0, title='Top-ranked factors')
sns.despine()
plt.tight_layout();
```



```
In [ ]:
```